

Deep Generative Model Tutorial

VAE & GAN

Taewook Nam
MLAI, KAIST
namsan@kaist.ac.kr

In this tutorial,

- Demonstrate **VAE & GAN** with toy data (MNIST)
- Describe brief outline of implementation
- You can go in detail with code : github.com/namsan96/1126_vaegan

code & slide :

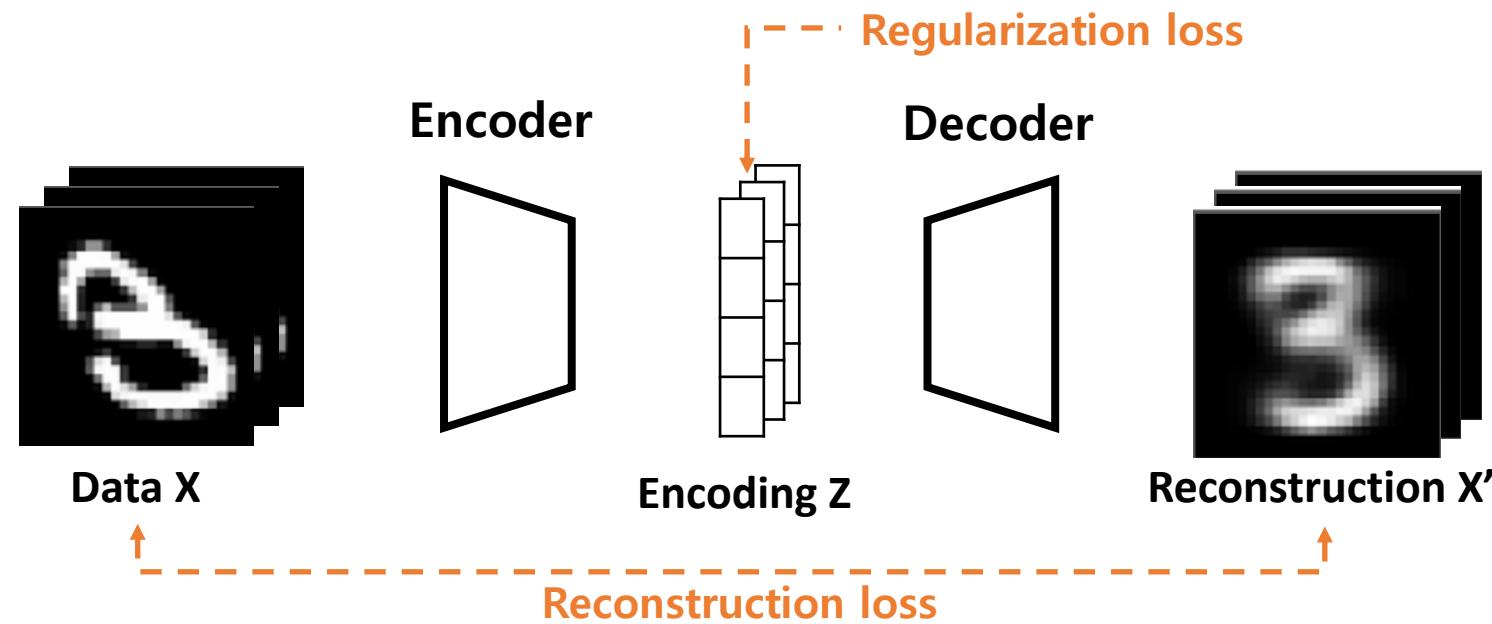
github.com/namsan96/1126_vaegan

Variational Auto-Encoder (VAE)

code & slide :
github.com/namsan96/1126_vaegan



VAE Overview



code & slide :
github.com/namsan96/1126_vaegan

Implementation TODO

- Encoder & Decoder Network
- Loss function
- Training Loop
- Visualization

code & slide :

github.com/namsan96/1126_vaegan

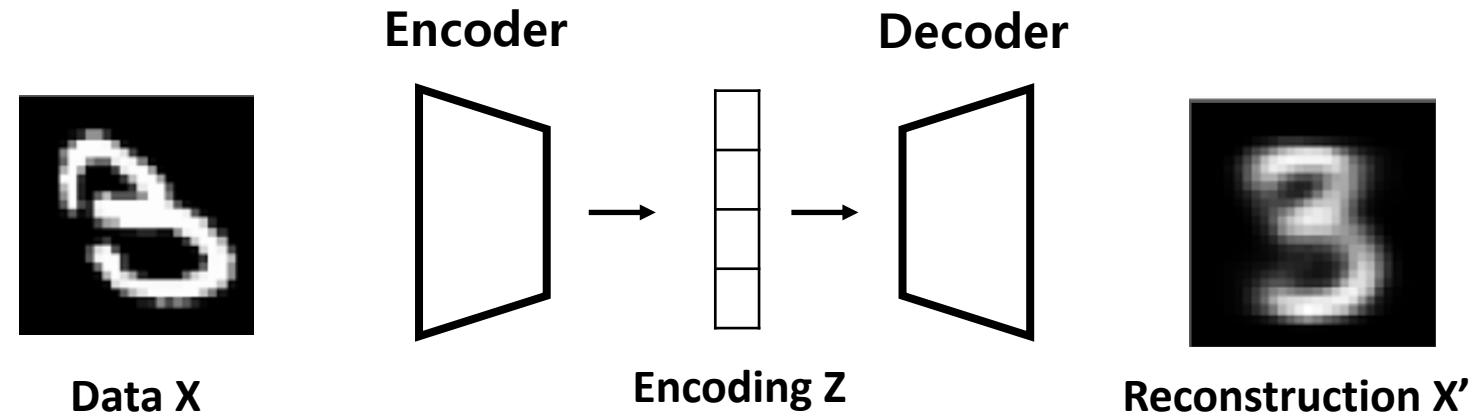
Implementation TODO

- **Encoder & Decoder Network**
- Loss function
- Training Loop
- Visualization

code & slide :

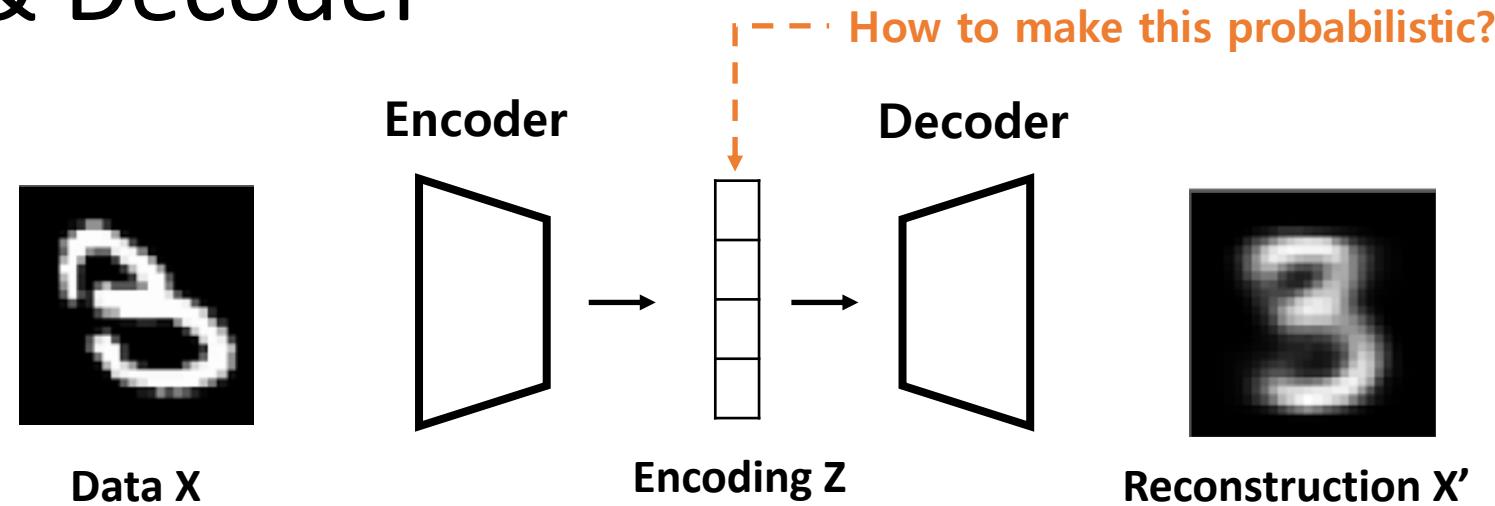
github.com/namsan96/1126_vaegan

Encoder & Decoder



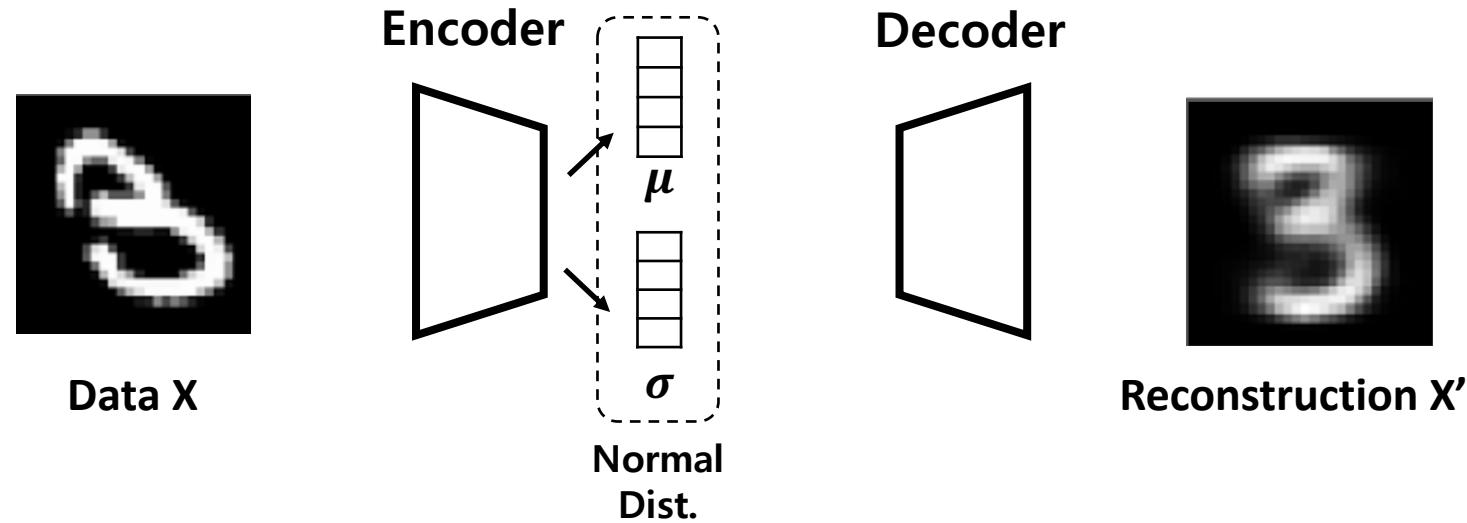
code & slide :
github.com/namsan96/1126_vaegan

Encoder & Decoder



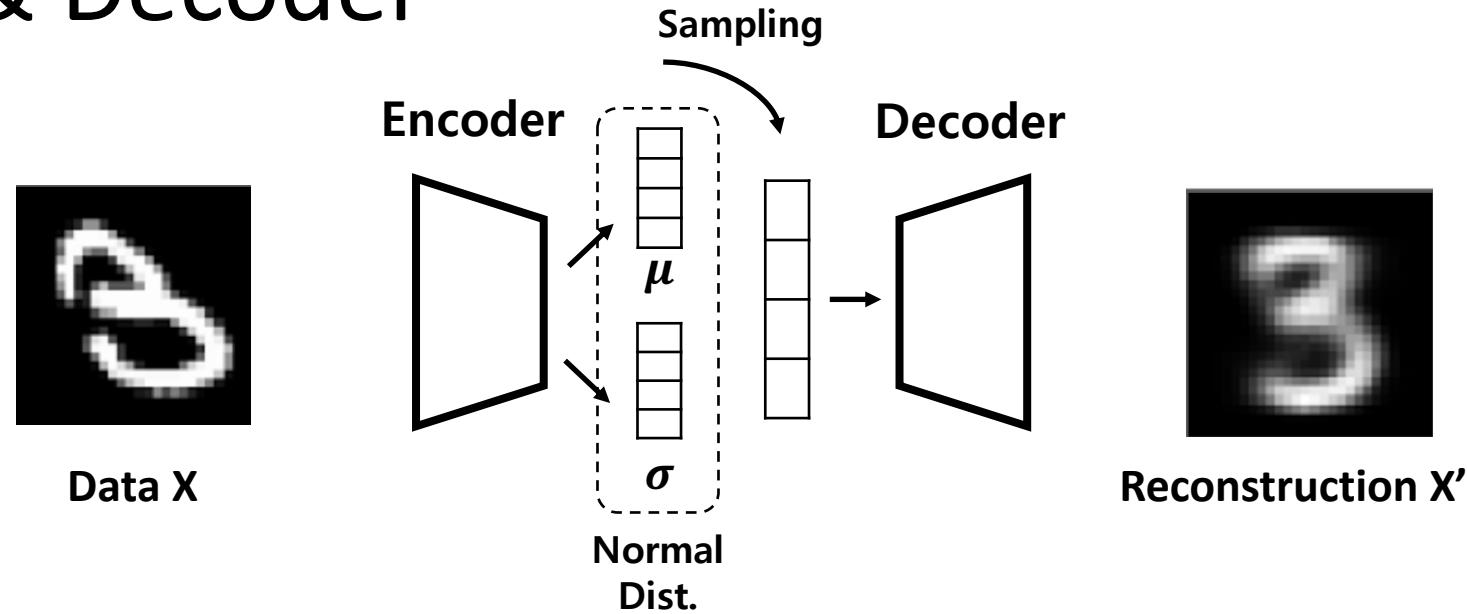
code & slide :
github.com/namsan96/1126_vaegan

Encoder & Decoder



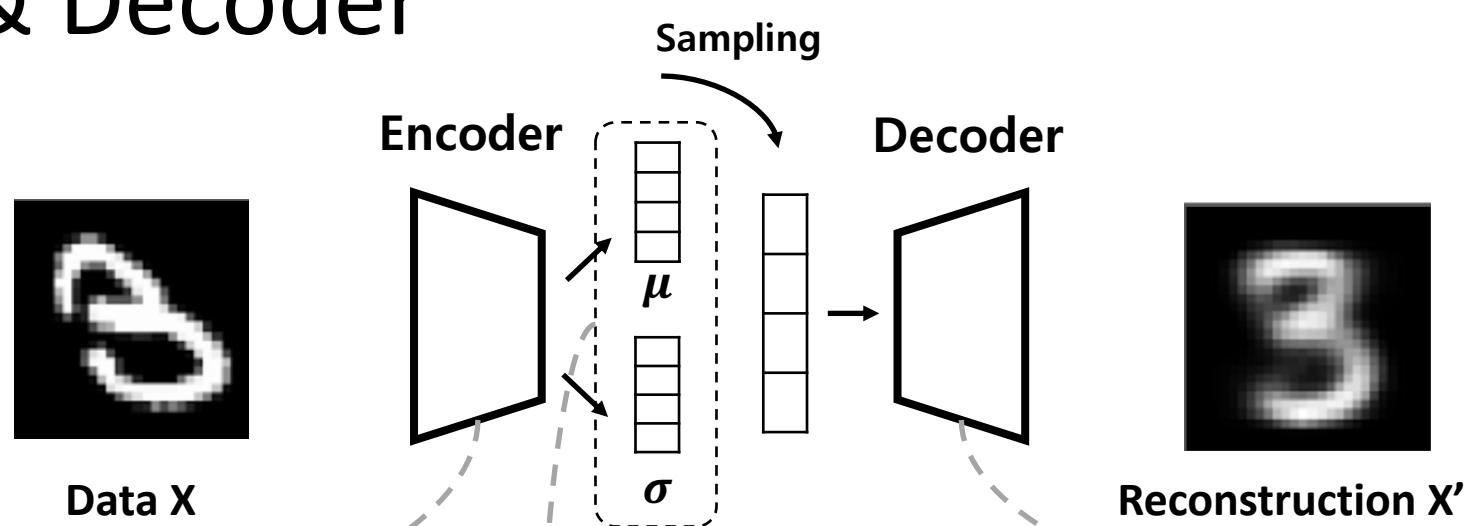
code & slide :
github.com/namsan96/1126_vaegan

Encoder & Decoder



code & slide :
github.com/namsan96/1126_vaegan

Encoder & Decoder



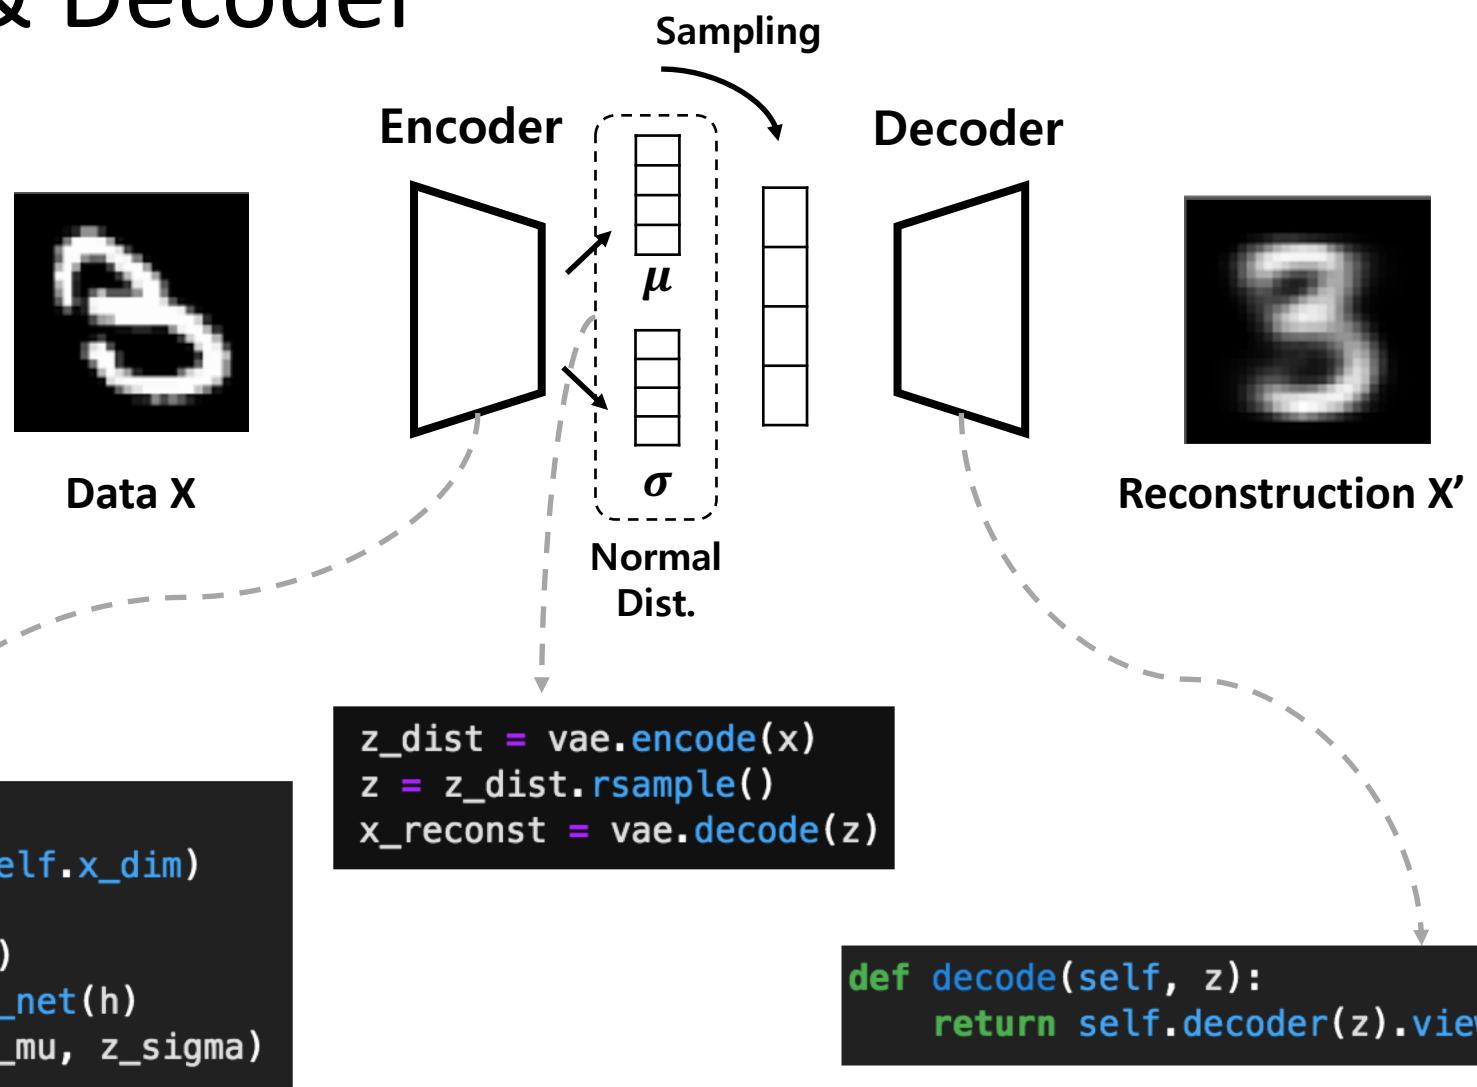
```
self.encoder = nn.Sequential(  
    nn.Linear(self.x_dim, hidden_dim),  
    nn.ReLU(),  
    nn.Linear(hidden_dim, hidden_dim),  
    nn.ReLU(),  
    nn.Linear(hidden_dim, hidden_dim),  
    nn.ReLU(),  
)
```

```
self.mu_net = nn.Sequential(  
    nn.Linear(hidden_dim, z_dim)  
)  
self.sigma_net = nn.Sequential(  
    nn.Linear(hidden_dim, z_dim),  
    nn.Softplus()  
)
```

```
self.decoder = nn.Sequential(  
    nn.Linear(z_dim, hidden_dim),  
    nn.ReLU(),  
    nn.Linear(hidden_dim, hidden_dim),  
    nn.ReLU(),  
    nn.Linear(hidden_dim, self.x_dim),  
    nn.Sigmoid(),  
)
```

code & slide :
github.com/namsan96/1126_vaegan

Encoder & Decoder



code & slide :
github.com/namsan96/1126_vaegan

Implementation TODO

- Encoder & Decoder Network
- **Loss function**
- Training Loop
- Visualization

code & slide :

github.com/namsan96/1126_vaegan

ELBO

$$\max_{\theta, \phi} \mathbb{E}_{q(\mathbf{z}|\mathbf{x}, \phi)} [\log p(\mathbf{x}|\mathbf{z}, \theta)] - KL[q(\mathbf{z}|\mathbf{x}, \phi) || p(\mathbf{z})]$$

$$\begin{aligned}\log p(\mathbf{x}|\mathbf{z}, \theta) &= \log \prod_{d=1}^{D_x} Ber(x_d|x'_d) \\ &= \sum_{d=1}^D \log(x'_d)^{x_d}(1-x'_d)^{x'_d} \\ &= \sum_{d=1}^D x_d \log x'_d + (1-x_d) \log(1-x'_d)\end{aligned}$$

$$\begin{aligned}KL[q(\mathbf{z}|\mathbf{x}, \phi) || p(\mathbf{z})] &= \sum_{d=1}^{D_z} KL[N(z_d|\mu_d, \sigma_d^2) || N(z_d|0,1)] \\ &= \sum_{d=1}^{D_z} \mu_d^2 + \sigma_d^2 - \log \sigma_d^2 - 1\end{aligned}$$

```
def compute_elbo(x, x_reconst, z_posterior, z_prior):
    log_likelihood = (x*torch.log(x_reconst + 1e-6) + (1-x)*torch.log(1 - x_reconst + 1e-6)).sum((1, 2, 3))
    kl = dist.kl_divergence(z_dist, z_prior).sum(1)
    return (log_likelihood - kl).mean()
```

code & slide :
github.com/namsan96/1126_vaegan

Implementation TODO

- Encoder & Decoder Network
- Loss function
- **Training Loop**
- Visualization

code & slide :

github.com/namsan96/1126_vaegan

Training Loop

```
for epoch_i in trange(n_epoch):
    for x, _ in loader:
        x = x.to(device)

        z_dist = vae.encode(x)
        z = z_dist.rsample()
        x_reconst = vae.decode(z)

        elbo = compute_elbo(x, x_reconst, z_dist, z_prior)
        loss = -elbo

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
```

code & slide :
github.com/namsan96/1126_vaegan

Implementation TODO

- Encoder & Decoder Network
- Loss function
- Training Loop
- **Visualization**

code & slide :

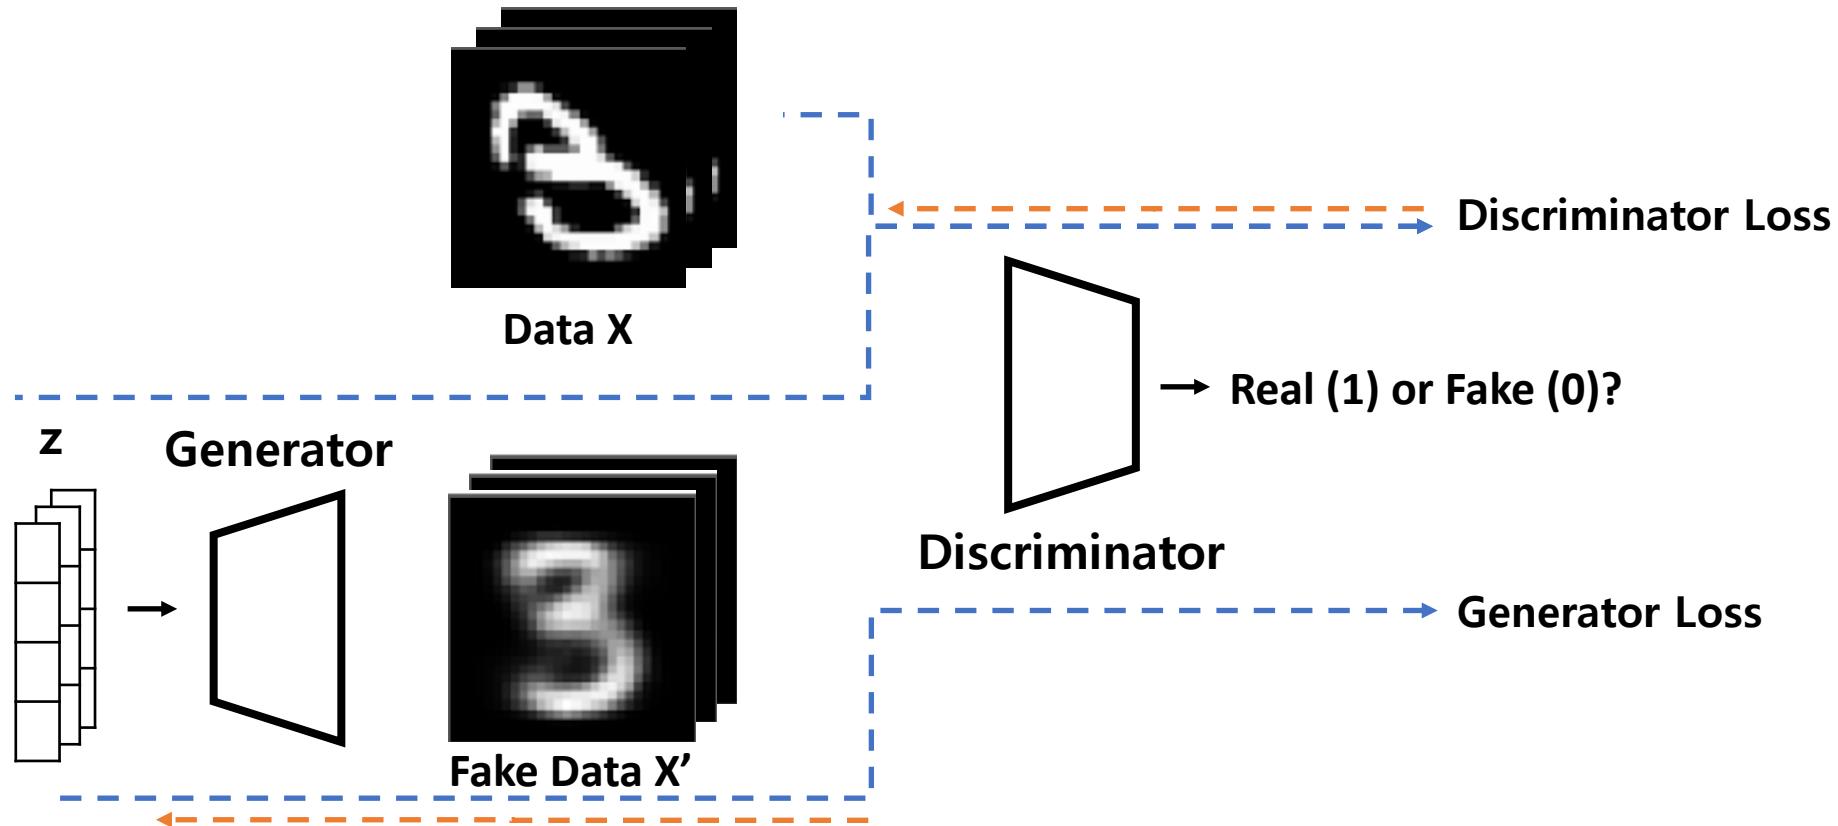
github.com/namsan96/1126_vaegan

Generative Adversarial Networks (GAN)

code & slide :
github.com/namsan96/1126_vaegan



GAN Overview



code & slide :
github.com/namsan96/1126_vaegan

Implementation TODO

- Generator & Discriminator Network
- Training Loop
- Visualization

code & slide :
github.com/namsan96/1126_vaegan

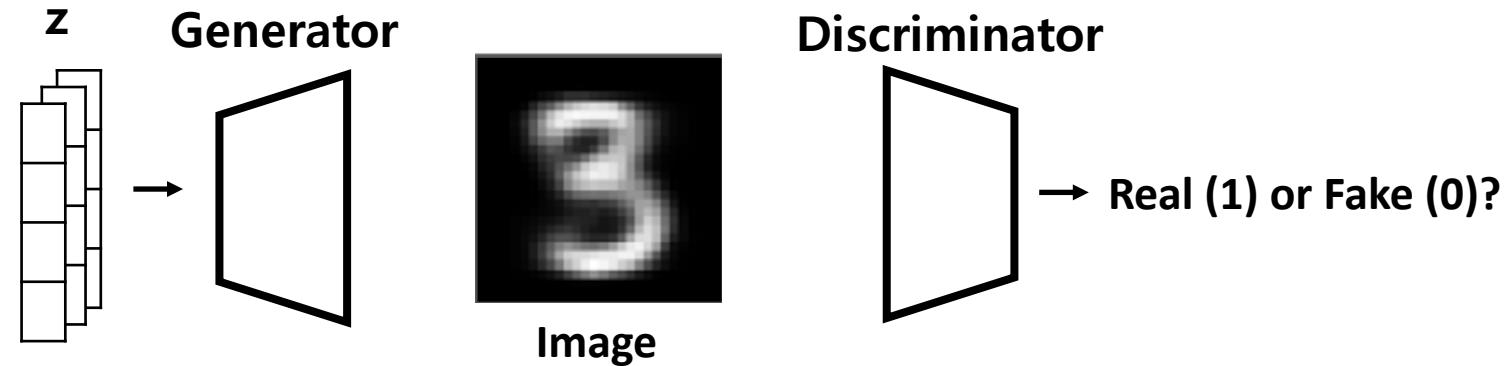
Implementation TODO

- **Generator & Discriminator Network**
- Training Loop
- Visualization

code & slide :

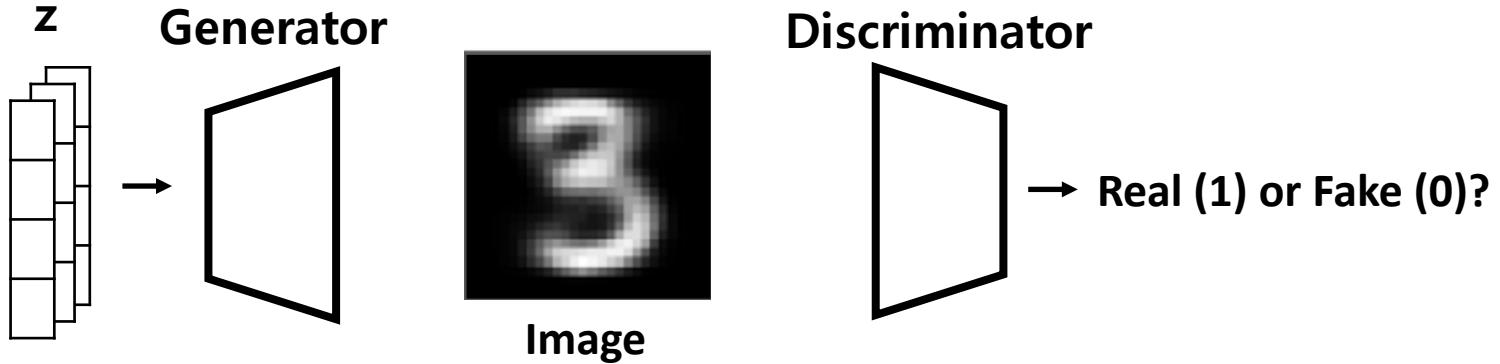
github.com/namsan96/1126_vaegan

Generator & Discriminator



code & slide :
github.com/namsan96/1126_vaegan

Generator & Discriminator



```
class G(nn.Module):
    def __init__(self, z_dim, x_shape, hidden_dim=500):
        super().__init__()

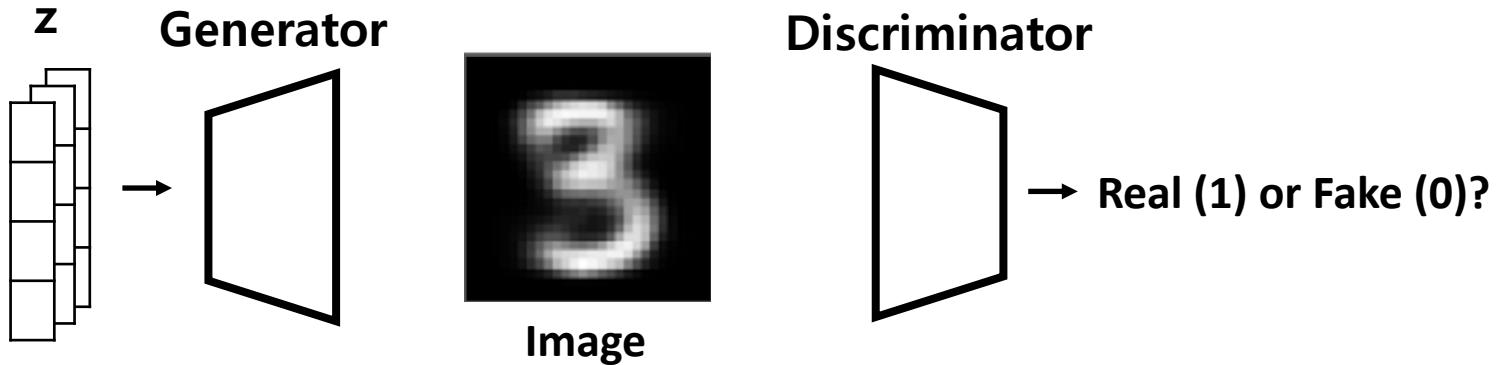
        self.x_shape = x_shape
        self.x_dim = np.prod(x_shape)
        self.z_dim = z_dim

        self.net = nn.Sequential(
            nn.Linear(z_dim, hidden_dim),
            nn.ReLU(),
            nn.Linear(hidden_dim, hidden_dim),
            nn.ReLU(),
            nn.Linear(hidden_dim, self.x_dim),
            nn.Sigmoid()
        )

    def forward(self, input):
        return self.net(input).view(len(input), *self.x_shape)
```

code & slide :
github.com/namsan96/1126_vaegan

Generator & Discriminator



```
class D(nn.Module):
    def __init__(self, x_shape, hidden_dim=500):
        super().__init__()

        self.x_dim = np.prod(x_shape)

        self.net = nn.Sequential(
            nn.Linear(self.x_dim, hidden_dim),
            nn.ReLU(),
            nn.Linear(hidden_dim, hidden_dim),
            nn.ReLU(),
            nn.Linear(hidden_dim, 1),
            nn.Sigmoid()
        )

    def forward(self, input):
        return self.net(input.view(len(input), self.x_dim)).squeeze(-1)
```

code & slide :

github.com/namsan96/1126_vaegan

Implementation TODO

- Generator & Discriminator Network
- **Training Loop**
- Visualization

code & slide :

github.com/namsan96/1126_vaegan

Training Loop

```
for epoch_i in range(n_epoch):  
    for real_x, _ in loader:
```

```
        d.zero_grad()  
  
        # discriminator for real data  
        real_x = real_x.to(device)  
        real_pred_y = d(real_x)  
        real_y = make_real_y(len(real_x)).to(device)  
  
        real_d_loss = criterion(real_pred_y, real_y)  
        real_d_loss.backward()  
  
        # discriminator for fake data  
        z = torch.randn(len(real_x), z_dim, device=device)  
        fake_x = g(z)  
        fake_pred_y = d(fake_x.detach())  
        fake_y = make_fake_y(len(real_x)).to(device)  
  
        fake_d_loss = criterion(fake_pred_y, fake_y)  
        fake_d_loss.backward()  
  
        d_optimizer.step()
```

```
        # generator  
        g.zero_grad()  
        fake_pred_y = d(fake_x).view(-1)  
        g_loss = criterion(fake_pred_y, real_y)  
        g_loss.backward()  
        g_optimizer.step()
```

Implementation TODO

- Generator & Discriminator Network
- Training Loop
- **Visualization**

code & slide :

github.com/namsan96/1126_vaegan

In this tutorial,

- Demonstrate **VAE & GAN** with toy data (MNIST)
- Describe brief outline of implementation
- You can go in detail with code : github.com/namsan96/1126_vaegan

code & slide :

github.com/namsan96/1126_vaegan

Q&A

code & slide :
github.com/namsan96/1126_vaegan



code & slide :
github.com/namsan96/1126_vaegan

